



One pane of glass for your BI + AI estate

A vendor-agnostic experience layer that hosts any BI tool and pairs it with the organisation's own AI — grounded, governed, built on platforms you already own.

Internal enabler · Path C

10 AI connector paths

\$0 software licence

v0.1 · phased rollout

THE PROBLEM

AI is trapped inside each BI tool

We run multiple BI tools and own multiple AI services — but every AI capability is locked to the tool that ships it.

The silo tax

- Power BI's AI only sees Power BI; Tableau's only Tableau.
- You can't bring your own approved model or governed Genie space.
- A different "ask the data" experience per tool.
- Governance fragmented — each tool mints its own tokens.

Why the timing is right

- We already own Databricks (Genie, Unity Catalog, Mosaic models & agents) + cloud LLMs.
- Those investments are under-leveraged — no single surface points them at what users look at.
- A thin layer turns existing spend into one governed, plain-language experience — no new licences.

WHAT PULSEPLAY IS

Not another BI tool. Not another chatbot.

A context-aware, vendor-neutral experience + governance layer that connects the BI surfaces and the AI brains we already run — so anyone can look at any dashboard and ask it questions in plain language, with answers grounded in governed data.

2**Independent axes**

Swap BI vendor and AI brain independently.

10**AI connectors**

One contract, ten backend paths.

\$0**New licence**

Internally built; reuse owned platforms.


THE DEFINING DESIGN

Two independent axes

Y — BI vendor (what you look at)


MOVING PARTS

Three tiers — the proxy is the only credential holder



REQUEST LIFECYCLE

How a question becomes a trusted answer



PLUGGABLE BRAINS

10 AI connector paths — all implemented

#	Connector	Type
1	Databricks Genie	NL → governed SQL
2	Azure OpenAI — chat	LLM
3	Azure OpenAI — analytics	LLM + schema
4	Bedrock — Retrieve&Generate	RAG
5	Bedrock — InvokeModel	direct model

#	Connector	Type
6	Mosaic Foundation Model	served endpoint
7	Supervisor — endpoint	multi-agent
8	Supervisor — local fan-out	N-space synth
9	Mosaic ResponsesAgent	agent endpoint
10	Power BI DAX	no-LLM · \$0 tokens

Why it matters: different questions warrant different brains — a governed Genie space for grounded SQL, a deterministic DAX path for zero-cost exact KPIs, a Foundation Model for narrative. Chosen per profile, changed without touching the product.

FEATURES

What it does today

Host any dashboard LIVE

Power BI real integration;
Tableau/Qlik/Looker embed via iframe.

AI Insights LIVE

Auto briefing: Headline, KPIs, Trends, Risks,
Actions + custom sections.

Ask Pulse LIVE

Plain-language Q&A grounded in the active
BI context.

Workbench LIVE

3 tabs + author-only templates (tabs,
landing, scope, sections).

Author controls LIVE

Markdown / SQL sections, presets
(SWOT/BCG), format & masking rules.

Governed by design LIVE

Allow-lists, server-side tokens,
masking/redaction, 4-status trust.

BENEFITS

Why the business should care

Benefit	What it means	Who feels it
No vendor lock-in	Change BI tool or AI model via config; experience stays constant	Architecture
Reuse existing spend	Genie, Foundation Models, UC, BI licences leveraged — not duplicated	Finance
One experience	Users learn one “ask the data” surface across every dashboard	End users
Central governance	Tokens, masking, allow-lists, audit in one proxy	Security
Grounded answers	From governed data, with an explicit trust status	Decision-makers
Efficiency	“Fewer tokens, better accuracy” — curbs inference spend vs a naive bot	Finance

“Look at anything. Ask in plain language. Trust the answer.”

COMPARATIVE ANALYSIS

By capability, not feature-for-feature

Capability	Native vendor AI	Genie direct	Build per tool	PulsePlay
Across multiple BI vendors	✗ single	✗ data-only	△ per app	✓
Swap the AI model	✗ vendor's	△ Genie	✓ if built	✓ 10 paths
Reasons about what you see	✓ in-tool	✗	△ wire it	✓
Central governance/tokens	✗ per tool	△ UC	△ build	✓ one proxy
Trust/grounding contract	△ varies	△ raw	△ build	✓ 4-status
New software licence	\$\$ per seat	\$ compute	\$\$ build+run	\$0 internal
Answer-quality maturity	✓ mature	✓ mature	△ depends	△ improving

Honest line: if you're single-vendor and the built-in AI suffices, native Copilot is mature and may be enough. PulsePlay's edge shows the moment you have more than one BI tool, want to choose the AI brain, or need uniform governance.

COSTING MODEL

No software licence — people + pass-through compute

1 · People (the real new cost)

- Engineering ownership (React app + proxy, new adapters/connectors)
- Platform/DevOps (deploy, monitor, secrets & rotation)
- BI / data admins (Genie spaces, UC policy, Power BI capacity)

2 · Platform (already owned)

- Databricks warehouse compute + Foundation Model serving
- LLM token inference (Azure OpenAI / Bedrock) — pay-per-use
- Power BI Premium · App hosting · Identity · Key Vault · SIEM

“No new software to buy. We fund a small support team and pay only for the compute & tokens we actually use, on platforms we already run.”

HONEST MATURITY

Solid foundation, phased rollout

Area	Status
2-axis architecture · 10 connectors · governance · Power BI embed tokens	SOLID & TESTED
Power BI adapter · AI Insights · Ask Pulse · Workbench · authoring · masking · theming	LIVE
Tableau / Qlik / Looker real SDK adapters	IFRAME TODAY · ROADMAPPED
Answer-quality scoring (ground-truth eval)	QUALITATIVE · ROADMAPPED
Streaming · full RAG/citations · multi-tenant · public-OSS	NOT IN CURRENT SCOPE

Structural correctness (routing, governance, section shape, UI) is proven by tests. **Answer correctness** is qualitative-review today — we present it as trust-tagged decision support, not an oracle.

CAVEATS

What NOT to commit to

Caveat	Say this
Vendor adapters	Power BI is deep; Tableau/Qlik/Looker embed today, deep SDKs roadmapped
Genie Agent Mode	UI-only in the API; we use the Foundation Model path instead
Power BI RLS	Requires Premium capacity
Masking	Defence-in-depth — Unity Catalog column masks are the real control
Answer quality	Validated for structure, not yet scored for factual correctness
Scope	Internal enabler — multi-tenant & public-OSS are deferred

Golden rule: commit to the [architecture](#), [governance](#), and the [Power BI + Genie experience](#). Frame everything else as [roadmap](#), not “done.”

CROSS-EXAMINATION CHEAT SHEET

If they ask... you say

"Why not just Power BI Copilot?"

It only works in Power BI with Microsoft's model. We run multiple tools and want our own approved brain, governed centrally.

"Are you building an LLM/agent?"

No — we orchestrate Genie, Foundation Models, agents we already own. PulsePlay is the thin experience + governance layer.

"Can we trust the AI for decisions?"

Grounded in governed data, every answer trust-tagged by our validator. Structure enforced today; quality scoring roadmapped. Decision-support with explicit trust.

"Do credentials/data leak?"

Credentials never touch the browser; the proxy mints scoped tokens. Sensitive fields masked before any model. UC is the boundary.

"What does it cost?"

No licence — internally built. A small support team + only the compute/tokens we use on platforms we already pay for.

"Is it production-ready?"

Architecture & backbone tested; Power BI + Genie is the hardening focus. Internal v0.x — solid base, phased rollout.

IN ONE LINE

PulsePlay turns siloed BI + AI investments into one governed, plain-language experience.

No new licences · honest about maturity · built to outlast any single vendor or model choice.