



EXECUTIVE BRIEFING

One pane of glass for your BI + AI estate

A vendor-agnostic experience layer that hosts any BI tool and pairs it with the organisation's own AI — grounded, governed, and built on platforms you already own.

How to use this briefing

This document is written so a non-technical sponsor can present it *and* defend it. Each section gives the executive message first; the tables and the “Cross-examination cheat sheet” (§12) give the technical backup to answer push-back without over-committing.

- | | |
|--------------------------------|----------------------------------|
| 01 Executive summary | 08 Comparative analysis |
| 02 The problem & why now | 09 Costing model |
| 03 What PulsePlay is | 10 Security & governance |
| 04 Architecture & moving parts | 11 Maturity & honest status |
| 05 The 10 AI connector paths | 12 Cross-examination cheat sheet |
| 06 Features (the product) | 13 Caveats — what NOT to commit |
| 07 Benefits | 14 Roadmap & glossary |

The one-paragraph version

The organisation already owns BI tools (Power BI, Tableau, Qlik, Looker) *and* AI services (Databricks Genie, Mosaic Foundation Models & agents, Azure OpenAI, Bedrock). They sit in silos. **PulsePlay is the thin layer that puts them behind one screen**: the user looks at any dashboard, and an AI assistant — wired to whichever brain the org chooses — reasons about exactly what they're looking at, with answers that are **grounded in governed data** and **gated by a trust contract** so nothing ungrounded is presented as fact. We don't build LLMs, agents, or data platforms; we orchestrate what exists and own the experience + governance layer.

2**Independent axes**

Swap the **BI vendor** (what you look at) and the **AI brain** (what answers) independently. No lock-in to either.

10**AI connector paths**

Genie, Foundation Models, Azure OpenAI, Bedrock, Supervisor agents, Power BI semantic-model — all behind one contract.

\$0**New software licence**

Inner-source & internally built. Cost is **people to support it** + platform you already pay for — not per-seat software.

Why this matters to the business

- **Consolidation, not another silo.** One assistant across heterogeneous BI tools and AI services — reusing existing investments rather than buying a new platform.
- **Future-proof by design.** The vendor and the AI brain are pluggable. When the org adopts a new BI tool or a better model, you change a setting — not the product.
- **Governed and grounded.** Embed tokens are issued server-side, sensitive values are masked/redacted before they reach any model, access is allow-listed, and a validator gates every AI artifact with an honest status (Verified / Grounded draft / Suggestion / Blocked).
- **Honest maturity.** The architecture and the AI connector backbone are solid and tested; Power BI is a real integration; Tableau/Qlik/Looker are iframe-embed today (real SDKs are roadmapped). We say so plainly so commitments stay safe.

The headline you can say out loud

"PulsePlay lets anyone ask questions in plain language about any dashboard we run, using our own AI, on our own data, with our own governance — without us buying or rebuilding anything."

Enterprises have accumulated multiple BI tools and, separately, multiple AI services. Each AI capability is trapped inside the tool that ships it.

The silo tax

- Power BI's AI only sees Power BI. Tableau's only sees Tableau. Qlik's only sees Qlik.
- Each vendor's assistant is tied to that vendor's model — you can't bring your own approved LLM or your governed Genie space.
- Users learn a different "ask the data" experience per tool; analysts rebuild the same prompt logic everywhere.
- Governance is fragmented: every tool mints its own tokens, applies its own (or no) masking.

Why the timing is right


- The org has already invested in Databricks (Genie, Unity Catalog, Mosaic Foundation Models & agents) and in cloud LLMs.
- Those investments are under-leveraged because there's no consistent surface that points them at what users actually look at.
- A thin orchestration layer turns existing spend into a single, governed, plain-language experience — fast, and without new licences.

The gap PulsePlay fills

Not "another BI tool" and not "another chatbot." The missing piece is a **context-aware, vendor-neutral experience + governance layer** that connects the BI surfaces and the AI brains the org already runs.

What PulsePlay is — the 2-axis design

PulsePlay's defining idea: two independent axes. **What you look at** (BI vendor) and **what answers you** (AI connector) are decoupled. Any combination is valid; switching either is independent of the other.



Any BI vendor (Y) × any AI connector (X). Switching one never forces a change to the other.

It is the host, not a guest


PulsePlay runs at the top level of a real modern browser and embeds vendor surfaces inside narrowed sandbox iframes that *it* controls. It is not trapped inside any vendor's sandbox, so modern web capabilities (streaming, exports, rich charts) are available.

It orchestrates; it doesn't rebuild

No new LLM, no new agent framework, no new data warehouse. PulsePlay calls the org's existing Genie spaces, Foundation Model endpoints, Supervisor agents and BI tools — and adds the experience, grounding and governance around them.

04 Architecture & moving parts

Three tiers: a browser experience layer, a connector-agnostic proxy (the security & routing backbone), and the platform services the org already owns. The proxy is the only place that holds credentials and mints tokens.



The browser never holds a credential. All token issuance, routing and policy live in the proxy.

Layer	Technology	Responsibility
Experience	React · Vite · TypeScript	Hosts BI vendor surfaces; AI Insights + Ask Pulse; validator & trust display; author Settings + governance UI; theming.
BI adapters (Y)	BIAdapter contract; powerbi-client SDK	One contract every vendor implements. Power BI is a real SDK adapter; others embed via iframe today.
Proxy (X)	Node · Express	The only credential holder. Routes to 10 AI backends; mints embed tokens; sanitises prompts; redacts secrets/PII; gates SQL; rate-limits; audits.
Databricks	Genie · Mosaic FM/agents · Unity Catalog · SQL Warehouse	The data + AI brains. Unity Catalog is the real data-access boundary (row/column security).
Visualisation	Power BI native canvas · Markdown · (ECharts/Vega-Lite roadmapped)	Renders the embedded dashboard + AI answers/sections.

The 10 AI connector paths

The proxy speaks one internal contract and translates it to whichever AI backend a profile selects. This is what makes the AI brain pluggable. All ten are implemented today.

#	Connector path	What it is	Notes
1	Databricks Genie	Natural-language → governed SQL over a Genie space	Primary grounded path
2	Azure OpenAI — chat	Conversational LLM	Bring-your-own approved model
3	Azure OpenAI — analytics	LLM with schema-aware analytics mode	
4	AWS Bedrock — Retrieve&Generate	RAG over a Bedrock knowledge base	
5	AWS Bedrock — InvokeModel	Direct model invocation	
6	Mosaic AI Foundation Model	Databricks-served model endpoint	Workaround for Genie Agent-Mode limits (see §13)
7	Supervisor — real endpoint	Multi-agent orchestration (LangGraph)	
8	Supervisor — local fan-out	Proxy fans out to N Genie spaces + synthesises	Cost scales $\sim(N+1)$ calls — see §13
9	Mosaic AI ResponsesAgent	Databricks ResponsesAgent endpoint	
10	Power BI semantic-model (DAX)	Deterministic DAX templates — no LLM	Zero token cost; fully deterministic answers

Why ten matters commercially

Different questions warrant different brains: a governed Genie space for grounded SQL, a deterministic DAX path for zero-cost exact KPIs, a Foundation Model for narrative. The org picks per profile — and can change its mind without touching the product.

Host any dashboard LIVE

Embed the vendor surface; Power BI is a real client integration with secure embed + a developer tools strip. Tableau/Qlik/Looker embed via iframe today.

AI Insights briefing LIVE

Auto-generated executive briefing — Headline, KPI snapshot, Trends, Risks, Recommended actions — plus author-defined sections. Staged reveal for perceived speed.

Ask Pulse (conversational) LIVE

Plain-language Q&A grounded in the active BI context, through whichever connector is selected.

Workbench (3 tabs) LIVE

AI Insights · Ask Pulse · Dashboard, one consistent shell. Author-only “Workbench templates” preset the tabs, landing surface, scope and section set.

Author authoring controls LIVE

Sections in markdown *or* warehouse-backed SQL config-items; number-format guidance; SWOT/BCG/RFM presets; domain guidance; knowledge packs.

Governed by design LIVE

Allow-listed vendors/profiles/origins; server-side embed tokens; PII redaction + field masking; the 4-status trust contract; capability-aware “requires X” affordances.

The accuracy contract — “no ungrounded artifacts”

Every AI artifact carries one of four honest statuses, set by PulsePlay's validator (never by the model):

VERIFIEDGROUNDING DRAFT

SUGGESTION BLOCKED. This is the difference between a demo chatbot and an enterprise tool: the system tells the user how much to trust each answer. (Today the validator enforces section *shape/structure*; factual-correctness scoring is a roadmap item — see §11.)

Benefit	What it means	Who feels it
No vendor lock-in	Change BI tool or AI model via configuration; the experience stays constant.	Architecture / procurement
Reuse existing spend	Genie, Foundation Models, UC, BI licences are leveraged — not duplicated.	Finance / platform owners
One experience, many tools	Users learn one “ask the data” surface across every dashboard.	End users / analysts
Governance in one place	Tokens, masking, allow-lists, audit centralised in the proxy.	Security / compliance
Grounded answers	Answers come from governed data with an explicit trust status.	Decision-makers
Fast to extend	New connector or adapter is a drop-in module, not a rewrite.	Engineering
Efficiency motto	“Fewer tokens, better accuracy” — grounded/staged approach curbs inference spend vs a naive chatbot.	Finance / sustainability

Marketing one-liners (for the flyer / intranet)

- “Look at anything. Ask in plain language. Trust the answer.”
- “Your BI tools. Your AI. Your data. One screen.”
- “Bring your own brain — PulsePlay is connector-agnostic and vendor-neutral.”

PulsePlay is not competing with a single product — it sits *above* the BI tools and *beside* the AI services. The honest comparison is by capability, not feature-for-feature.

Capability	Native vendor AI PBI Copilot · Tableau Pulse · Qlik	Databricks Genie used directly	Build a custom AI app per tool	PulsePlay
Works across multiple BI vendors	✗ single-vendor	✗ data-only	△ per app	✓ vendor-neutral
Bring-your-own / swap the AI model	✗ vendor's model	△ Genie only	✓ if you build it	✓ 10 connector paths
Reasons about what you're looking at	✓ within its tool	✗	△ you wire it	✓ context-aware
Centralised governance / tokens / masking	✗ per tool	△ UC only	△ you build it	✓ one proxy
Trust/grounding contract on answers	△ varies	△ raw	△ you build it	✓ 4-status validator
New software licence cost	\$\$ per seat/capacity	\$ compute	\$\$ build + run	\$0 licence (internal)
Answer quality maturity	✓ mature	✓ mature	△ depends	△ improving (see \$11)
Time-to-value vs build	✓ instant (if single-tool)	✓ instant (data)	✗ months per tool	✓ one layer, all tools

Where native vendor AI still wins

If the org is effectively single-vendor and that vendor's built-in AI meets the need, native Copilot is mature and may be sufficient. PulsePlay's edge appears the moment you have **more than one** BI tool, or want to **choose the AI brain**, or need **uniform governance**.

Honest positioning vs ThoughtSpot

The closest philosophical peer exposes one vendor's data to any LLM. PulsePlay does the inverse + more: it hosts *any* BI tool in the browser under one assistant with *any* connector. Complementary, not the same category.

Because this is internally built and inner-sourced, there is **no per-seat application licence**. Cost has two real buckets: **people to support it**, and **platform services you already pay for** (PulsePlay rides on them; it does not add a new platform).

1 · People (the real new cost)

- Engineering ownership: maintain/extend the React app + proxy, add adapters/connectors.
- Platform/DevOps: deploy + monitor the proxy and app, manage secrets & rotation.
- BI / data platform admins: manage Genie spaces, Unity Catalog policy, Power BI capacity.
- This is the headcount commitment to be explicit about — the application itself is free of licence.

2 · Platform (pass-through, mostly already owned)

- Databricks SQL Warehouse compute (for the deterministic SQL/KPI path).
- Foundation Model serving + Azure OpenAI / Bedrock **token inference** — pay-per-use.
- Power BI Premium capacity (already owned for embedding + RLS).
- App hosting (Databricks App and/or Azure App Service tier).
- Identity, Key Vault, SIEM — existing enterprise services.

Cost lever	Driver	How PulsePlay keeps it down
LLM inference (tokens)	Insights + chat volume	Grounded + staged prompting; deterministic DAX path for exact KPIs costs zero tokens ; "fewer tokens, better accuracy" design.
Warehouse compute	SQL sections + Genie queries	Result caching; warehouse pre-warm/keep-alive to avoid repeated cold starts.
Hosting	Concurrent users	Lightweight React SPA + Node proxy; scales on standard App Service tiers.

The cost message in one line

"No new software to buy. We fund a small support team and pay only for the compute/tokens we actually use on platforms we already run."

Control	How it works
Credentials never in the browser	All embed tokens (Power BI, etc.) and AI credentials are issued/held server-side in the proxy. The browser bundle carries no secret.
Identity	Org IdP (SSO + MFA). No local accounts; PulsePlay does not own the user store.
Data access boundary	Unity Catalog row/column security is the real boundary. Power BI RLS is derived from verified identity claims (requires Premium — see §13).
Masking & PII redaction	Author-declared field masking (redact / last-4 / hide) + pattern-based PII redaction strip sensitive values <i>before</i> they reach any model, and from rendered cells. Defence-in-depth, not the primary control — Unity Catalog column masks are.
Allow-listing	Permitted BI vendors, AI profiles, embedding origins, Power BI workspaces, and AD-group access are all allow-listed; fail-closed when the policy can't be loaded.
Prompt & SQL safety	Prompt sanitisation + injection-keyword stripping; read-only SQL with a DML blocklist; rate limiting; full audit log with request-ID correlation.

Be precise on this in the room

PulsePlay's masking/redaction is **defence-in-depth**. The authoritative data-security control is **Unity Catalog column masks + RLS** at the data layer. Don't position PulsePlay's masking as the primary guarantee.

11 Maturity & honest status

10/10

Connector paths

All AI backend paths implemented + tested.

1 real

BI vendor (Power BI)

Tableau/Qlik/Looker embed via iframe today.

green

Test suites

Proxy + playground suites pass (structural correctness).

The architecture and the AI backbone are solid and well-tested. What's **structurally** proven (routing, governance, section shape, UI) is distinct from what's **qualitatively** proven (answer correctness). We are honest about the line:

Area	Status
2-axis architecture, 10 connector paths, governance/allow-list, embed-token issuance (Power BI)	SOLID & TESTED
Power BI real adapter; AI Insights + Ask Pulse + Workbench; author authoring; masking; theming	LIVE
Tableau / Qlik / Looker real SDK adapters	IFRAME TODAY · SDKS ROADMAPMED
Answer-quality evaluation (ground-truth scoring)	QUALITATIVE ONLY · EVAL SUITE ROADMAPMED
Token streaming, full RAG/citations, multi-tenant isolation, public-OSS	NOT IN CURRENT SCOPE

Likely technical push-backs and the safe, true answer. Use these verbatim.

Q “Why build this instead of just using Power BI Copilot / Tableau Pulse?”

A Those only work inside their own tool and their own model. We run more than one BI tool and we want to choose our own approved AI brain and govern it centrally. PulsePlay is the one layer that spans all of them — without per-tool rebuilds.

Q “Are you building your own LLM / agent / data platform?”

A No. We orchestrate Databricks Genie, Mosaic Foundation Models & agents, Azure OpenAI and Bedrock that we already own. PulsePlay is the experience + governance layer, deliberately thin.

Q “Is the AI accurate? Can we trust it for decisions?”

A Answers are grounded in governed data and every artifact carries a trust status (Verified / Grounded draft / Suggestion / Blocked) set by our validator, not the model. Today we enforce structure; ground-truth quality scoring is a roadmap item — so we present it as decision-support with explicit trust, not an oracle.

Q “Do credentials or data leak to the browser or to the model?”

A Credentials never touch the browser — the proxy holds them and mints scoped tokens server-side. Sensitive fields are masked/redacted before reaching any model. Unity Catalog row/column security is the authoritative boundary.

Q “You say Tableau/Qlik/Looker — is that real today?”

A Power BI is a real SDK integration. Tableau/Qlik/Looker embed via iframe today (they render and are governed, but without the deep event bridge); their real SDK adapters are roadmapped. We won't claim deep integration before it ships.

Q “What does it cost?”

A No application licence — it's internally built. We fund a small support/engineering team and pay only for the compute and tokens we actually use, on platforms (Databricks, Power BI, cloud LLMs) we already pay for.

Q “Can it scale / is it production-ready?”

A The architecture and backbone are tested and the Power BI + Genie path is the production-hardening focus. It's an internal v0.x: solid foundation, phased rollout. Multi-tenant isolation and public-OSS are explicitly out of current scope.

Q “Why not a single chatbot over a data lake?”

A A chatbot doesn't know what the user is looking at, doesn't host the governed dashboards, and doesn't enforce a trust contract. PulsePlay is context-aware, vendor-neutral and governed.

Read this before you promise anything. These are real, known limits — naming them protects the program.

Caveat	Don't say → Do say
Vendor adapters	Don't: "full Tableau/Qlik/Looker integration." → Do: "Power BI is deep; others embed today, deep SDKs are roadmapped."
Genie Agent Mode	The public Genie API ignores the deep-research flag (UI-only). We use the Foundation Model path as the workaround — don't promise Genie "agent mode" via API.
Genie messages immutable	One request = one message; multi-section briefings allocate multiple messages under a conversation. A technical constraint, handled — just don't design around editing a message in place.
Power BI RLS	Row-level security through Power BI embedding requires Premium capacity . Pro/Free can embed but not enforce RLS via the token.
Masking ≠ security guarantee	It's defence-in-depth. Unity Catalog column masks are the authoritative control.
Answer quality	Validated for structure, not yet scored for factual correctness. Position as trust-tagged decision support.
Supervisor fan-out cost	Querying N spaces costs $\sim(N+1)$ model calls per question — use deliberately, it scales linearly.
Power BI Q&A surface	Runs Microsoft's NLP in Microsoft's tenant (we only mint the token); it's in defer-mode toward its 2026 EOL — steer new work to Ask Pulse.
Scope	Internal-org enabler (Path C). Multi-tenant isolation, SBOM signing, conformance, public docs and full external compliance are deferred — not a public SaaS.

The golden rule for commitments

Commit to the **architecture, the governance model, and the Power BI + Genie experience**. Frame everything else (other vendors' deep SDKs, answer-quality scoring, streaming, multi-tenant) as **roadmap**, not "done."

Phased direction

- **Now (v0.x):** Harden the Power BI + Genie experience; author tooling; governance; dark theme; masking enforcement.
- **Next:** Real Tableau/Qlik/Looker SDK adapters; token streaming; answer-quality eval suite.
- **Later:** Full governed retrieval (RAG + citations); broader knowledge packs.
- **Deferred (separate decision):** Multi-tenant isolation, public-OSS readiness.

Glossary (for non-technical readers)

- **Connector / brain:** the AI service that answers (Genie, a model, an agent).
- **Adapter:** the code that embeds a specific BI vendor.
- **Genie:** Databricks' natural-language-to-SQL over governed data.
- **Unity Catalog (UC):** Databricks' data governance — the security boundary.
- **RLS/OLS:** row/object-level security — who can see which rows/fields.
- **Grounded:** the answer is derived from real, governed data, not invented.
- **Proxy:** the server that holds credentials, routes AI calls, enforces policy.

What to leave the room believing

PulsePlay turns existing, siloed BI + AI investments into one governed, plain-language experience — no new licences, honest about maturity, and built to outlast any single vendor or model choice.